

# Patagonia parrots density analysis

*Francisco Denes and Peter Solymos*

*June 8, 2018*

## Motivation

We want to model the density of the austral parakeet *Enicognathus ferrugineus* and the slender-billed parakeet *Enicognathus leptorhynchus* from count data obtained with road transect surveys in Patagonia. Because these parrots are gregarious, we want to assess whether group size and size of groups vary across habitats (classified as ‘urban’, ‘agropastoral’ and ‘other’ (i.e. various natural forest formations), and between breeding (Nov-Dec) and non-breeding seasons (Jan-Oct). We use distance sampling methods to account for detectability of parrots.

These animals tend to form smaller or larger groups, but groups of size 2 are also often observed as a result of mating pairs, as shown in Figure 1.




Figure 1: *Enicognathus ferrugineus* count frequencies

# Austral parakeet *Enicognathus ferrugineus*

## Estimating effective detection radius (EDR)

We fit distance sampling models with detection predictor variables, including the average group size, the number of groups that were observed (to account for the possibility that grouping behaviour influences detectability), and also habitat type. We used forward-stepwise model selection, starting with single covariate models and eliminating covariates that do not improve model parsimony (i.e. result in  $\Delta AIC < 2$ ) in relation to the null model.

Table 1: *E. ferrugineus* EDR models AIC

	df	AIC	$\Delta AIC$
EDR.habitatype	3	1434.532	0.00
EDR.null	1	1438.166	3.63
EDR.avggrousize	2	1439.144	4.61
EDR.numbergroups	2	1440.166	5.63

The model (EDR.habitatype) has the lowest AIC (Table 1), indicating that habitat type affects the effective detection radius (EDR) of *Enicognathus ferrugineus*. Specifically, detection radius is higher in agropastoral than urban and other habitats (Table 2). This means that detectability is higher in agropastoral habitats, and also that the area surveyed in a given sample unit is larger if the habitat therein is agropastoral vs. urban or other, presumably because it is possible to see further in pastures and planted fields than in forest or urban environments.

Table 2: *E. ferrugineus* top-ranked EDR model estimates

	Estimate	Std. Error	z value	$\Pr(> z )$
log.tau_(Intercept)	4.589	0.038	119.365	0.000
log.tau_Urban	0.057	0.058	0.988	0.323
log.tau_Agropastoral	0.263	0.103	2.556	0.011

Table 3: *E. ferrugineus* habitat-specific EDR (m)

Habitat	EDR
Other	98.36847
Urban	104.18432
Agropastoral	127.92056

## Models for number of groups

We then model the number of groups as a function of covariates. The model for number of groups is  $G_i \sim \text{Poisson}(D_i A_i)$ , where  $D_i$  = covariates and  $A_i$  = area sampled in site  $i$ .  $A_i$  is calculated using the habitat-specific estimated EDR as  $A_i = L_i * (2EDR_i)$ , where  $L$  is the length of each site (i.e. the transect).  $A_i$  is added to the model as an offset.

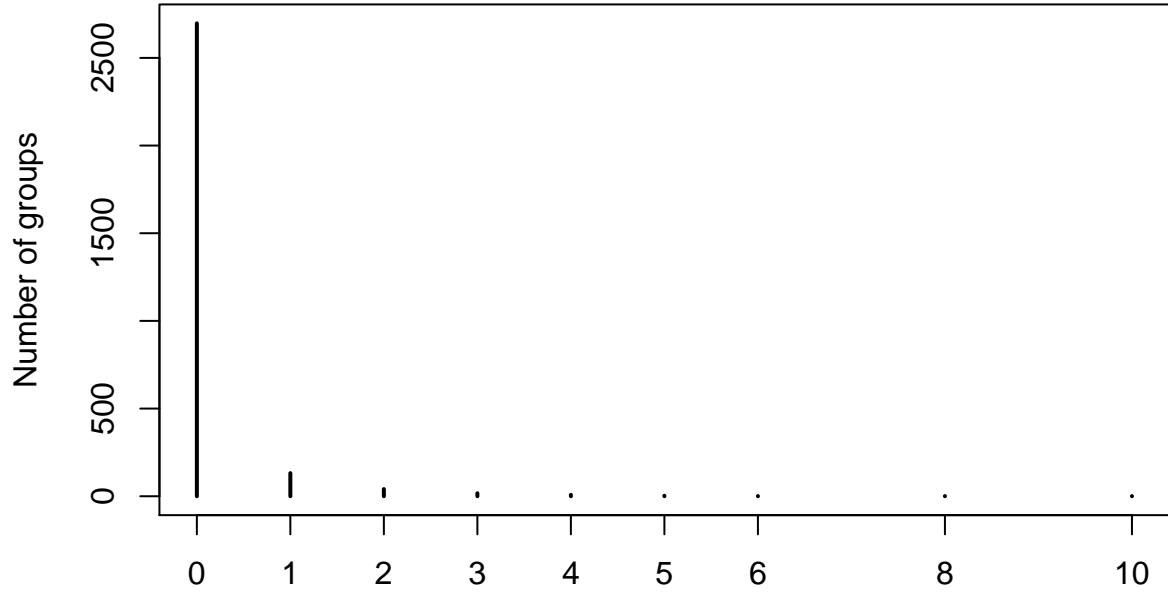


Figure 2: *Enicognathus ferrugineus* group numbers

### Model selection

For number of groups, we use a stage-wise selection procedure. First, we build a first set of models to evaluate the effect of covariates related to habitat (habitat type and elevation).

Table 4: *E. ferrugineus* number of group models

	df	AIC	dAIC
ngroup.hab.ele2	5	2082.963	0.00
ngroup.hab.ele	4	2114.437	31.47
ngroup.hab	3	2133.717	50.75
ngroup.ele2	3	2275.565	192.60
ngroup.ele	2	2323.806	240.84

The model with both habitat type and elevation (quadratic effect) has the lowest AIC (Table 4), indicating that both covariates affect the number of groups:

Table 5: *E. ferrugineus* ‘ngroup.hab.ele2’ model estimates

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-3.108	0.233	-13.356	0

	Estimate	Std. Error	z value	Pr(> z )
habitatOther	1.507	0.202	7.455	0
habitatUrban	2.390	0.202	11.823	0
elevation	-0.004	0.001	-7.032	0
I(elevation^2)	0.000	0.000	5.947	0

Table 6: Deviance partitioning of ‘ngroup.hab.ele2’ model for *E. ferrugineous*

	Df	Deviance	Resid. Df	Resid. Dev
NULL	NA	NA	2900	1865.338
habitat	2	207.965	2898	1657.373
elevation	1	21.279	2897	1636.094
I(elevation^2)	1	33.475	2896	1602.619

We then proceed by adding within-year temporal covariates (breeding/non-breeding season and Julian date) and their interactions with habitat.

Table 7: *E. ferrugineous* number of group models (within-year temporal predictors) AIC

	df	AIC	dAIC
ngroup.habXseason.ele2	8	2065.594	0.00
ngroup.hab.season.ele2	6	2067.555	1.96
ngroup.habXjdate.ele2	8	2081.546	15.95
ngroup.hab.ele2	5	2082.963	17.37
ngroup.hab.jdate.ele2	6	2084.491	18.90

The model with the season\*habitat interaction (Table 8) is equally parsimonious with the model with only the additive effects (Table 9). Given these are nested models, this indicates the interaction only marginally improves the model fit, so will continue with model with season and no interaction:

Table 8: *E. ferrugineous* ‘ngroup.habXseason.ele2’ interaction model estimates

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-18.187	558.012	-0.033	0.974
elevation	-0.004	0.001	-6.547	0.000
I(elevation^2)	0.000	0.000	5.500	0.000
habitatOther	15.797	558.012	0.028	0.977
habitatUrban	2.008	908.375	0.002	0.998
seasonnon-breeding	15.186	558.012	0.027	0.978
habitatOther:seasonnon-breeding	-14.357	558.012	-0.026	0.979
habitatUrban:seasonnon-breeding	0.266	908.375	0.000	1.000

Table 9: *E. ferrugineous* ‘ngroup.hab.season.ele2’ additive model estimates

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-4.250	0.412	-10.318	0.000
habitatOther	1.510	0.204	7.411	0.000
habitatUrban	2.314	0.203	11.424	0.000
elevation	-0.004	0.001	-6.408	0.000
I(elevation^2)	0.000	0.000	5.365	0.000
seasonnon-breeding	1.175	0.340	3.451	0.001

Finally, we assess year effects by adding a year covariate (2013-2016).

Table 10: *E. ferrugineous* number of group models (year predictor) AIC table

	df	AIC	dAIC
ngroup.hab.season.ele2.year	9	2022.265	0.00
ngroup.hab.season.ele2	6	2067.555	45.29

The model with lowest AIC indicates that the number of groups is affected by habitat type, elevation, season (breeding/non breeding) and year.

Table 11: *E. ferrugineous* top-ranked (ngroup.hab.season.ele2.year) model estimates

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-4.716	0.417	-11.311	0.000
habitatOther	1.739	0.206	8.434	0.000
habitatUrban	2.460	0.204	12.075	0.000
elevation	-0.003	0.001	-5.712	0.000
I(elevation^2)	0.000	0.000	4.588	0.000
seasonnon-breeding	1.022	0.348	2.939	0.003
as.factor(year)2014	0.092	0.168	0.549	0.583
as.factor(year)2015	0.918	0.161	5.688	0.000
as.factor(year)2016	0.612	0.205	2.979	0.003

Table 12: Deviance partitioning of top-ranked ‘ngroup.hab.season.ele2.year’ model for *E. ferrugineous*

	Df	Deviance	Resid. Df	Resid. Dev
NULL	NA	NA	2900	1865.338
habitat	2	207.965	2898	1657.373
elevation	1	21.279	2897	1636.094
I(elevation^2)	1	33.475	2896	1602.619
season	1	17.408	2895	1585.211
as.factor(year)	3	51.290	2892	1533.921

## Models for group size

The count distribution is characterized with a spike at 2 (Figure 1), and by the absence of 0s due to group size being conditional on having  $>0$  birds to consider it a group.

We start developing a general V-Inflated Poisson (VIP) model (“V” stands for “variable”, in reference to the “Z” representing 0 in a zero-inflated Poisson model, or ZIP), then we add the  $>0$  condition.

R functions for the VIP model are presented at the end of this document, including simulations to check the estimating procedure.

### Maximum likelihood

Let  $Y$  be a random variable, and  $y$  are observations,  $V$  is the count value that has some extra probability mass ( $V = 0$  is the ZIP model),  $f(y; \lambda)$  is the Poisson density ( $f(y; \lambda) = e^{-\lambda} \frac{\lambda^y}{y!}$ ).

The V-Inflated density can be written as  $P(Y = y) = \phi I(Y = V) + (1 - \phi)f(y; \lambda)$  which is  $\phi + (1 - \phi)f(V; \lambda)$  when  $Y = V$  and  $(1 - \phi)f(y; \lambda)$  otherwise.

We define the extra probability mass at  $V = 2$  to account for the group size peak in pairs. The model is also truncated at 0, because there are no groups with 0 individuals. We include habitat type as a covariate for the count (Poisson) component, and compare models with and without season (breeding/non-breeding) as a covariate for the V-inflation probability.

### Goodness-of-fit

Model goodness-of-fit can be evaluated visually, inspecting the proportions of fitted values against the count distribution. Deviation from the 1:1 line can be used as a goodness-of-fit metric, with the better model showing smaller deviation.




Figure 3: *Enicognathus ferrugineus* group size model goodness-of-fit (red: null model; blue: habitat.season model). Better model is closer to 1:1 line.

The deviation of null model is 1.38, and of the VIP.habitat.season model is 0.74.

### Confidence intervals for model estimates

Estimated model coefficients (Table 13) indicate that group sizes are larger in urban ( $\beta = 1.739$ ) and agropastoral ( $\beta = 0.614$ ) habitats when compared to other natural habitats. Moreover, the probability of there being extra pairs (i.e. the V-inflation probability,  $\phi$ ) is smaller in the non-breeding season ( $\beta = -2.326$ ).

Table 13: *E. ferrugineus* group size model estimates

	Estimate	Std. Error	z value	Pr(> z )
P_(Intercept)	1.007	0.019	53.262	0.000
P_Urban	1.739	0.023	75.931	0.000
P_Agropastoral	0.614	0.036	16.880	0.000
V_(Intercept)	0.223	0.671	0.333	0.740
V_seasonnon-breeding	-2.326	0.697	-3.337	0.001

Confidence intervals (CI) based on estimated standard errors can be obtained (Table 14):

Table 14: *E. ferrugineous* group size model CI

	2.5%	97.5%
P_(Intercept)	0.970	1.044
P_Urban	1.695	1.784
P_Agropastoral	0.543	0.685
V_(Intercept)	-1.092	1.538
V_seasonnon-breeding	-3.691	-0.960

Alternatively, we can estimate confidence intervals based on quantiles using bootstrap samples (with n=250) for the estimated coefficients (Table 15).

Table 15: Estimated coefficients and 95% CI based on bootstrap sample (n= 250) quantiles for *E. ferrugineous* group size models

	Estimate	2.5%	97.5%
P_(Intercept)	1.007	0.694	0.694
P_Urban	1.739	2.150	2.150
P_Agropastoral	0.614	0.565	0.565
V_(Intercept)	0.223	0.691	0.691
V_seasonnon-breeding	-2.326	-2.850	-2.850

## Density Predictions

The mean expected abundance ( $E$ ) can be predicted for each surveyed site  $i$  as  $E_i = (\phi_i 2 + (1 - \phi_i) \mu_i) N_i$ , where  $N$  is the mean number of groups,  $\mu$  is the mean group size and  $\phi$  is the V(=2)-inflation probability (i.e. the probability that a given group will be a pair).

We can then obtain the expected density stratified by habitat type, year and season, as  $D = \frac{\sum E}{\sum A}$ , where  $A$  is the area surveyed.

Table 16: *E. ferrugineous* mean density predictions for each habitat type, year and season

Habitat	Year	Season	Density (individuals/ha)	Density (2.5% CI)	Density (97.5% CI)
Agropastoral	2013	breeding	0.003	0.001	0.007
Other	2013	breeding	0.003	0.001	0.008
Urban	2013	breeding	0.009	0.003	0.021
Agropastoral	2015	breeding	0.001	0.000	0.003
Other	2015	breeding	0.003	0.001	0.008
Urban	2015	breeding	0.006	0.002	0.013
Agropastoral	2013	non-breeding	0.009	0.005	0.014
Other	2013	non-breeding	0.020	0.015	0.027
Urban	2013	non-breeding	0.135	0.094	0.186
Agropastoral	2014	non-breeding	0.011	0.007	0.016
Other	2014	non-breeding	0.019	0.014	0.025
Urban	2014	non-breeding	0.225	0.176	0.288
Agropastoral	2015	non-breeding	0.007	0.005	0.011
Other	2015	non-breeding	0.011	0.009	0.015
Urban	2015	non-breeding	0.089	0.069	0.114



Habitat	Year	Season	Density (individuals/ha)	Density (2.5% CI)	Density (97.5% CI)
Agropastoral	2016	non-breeding	0.003	0.002	0.005
Other	2016	non-breeding	0.009	0.006	0.012
Urban	2016	non-breeding	0.035	0.025	0.049

## Slender-billed parakeet *Enicognathus leptorhynchus*

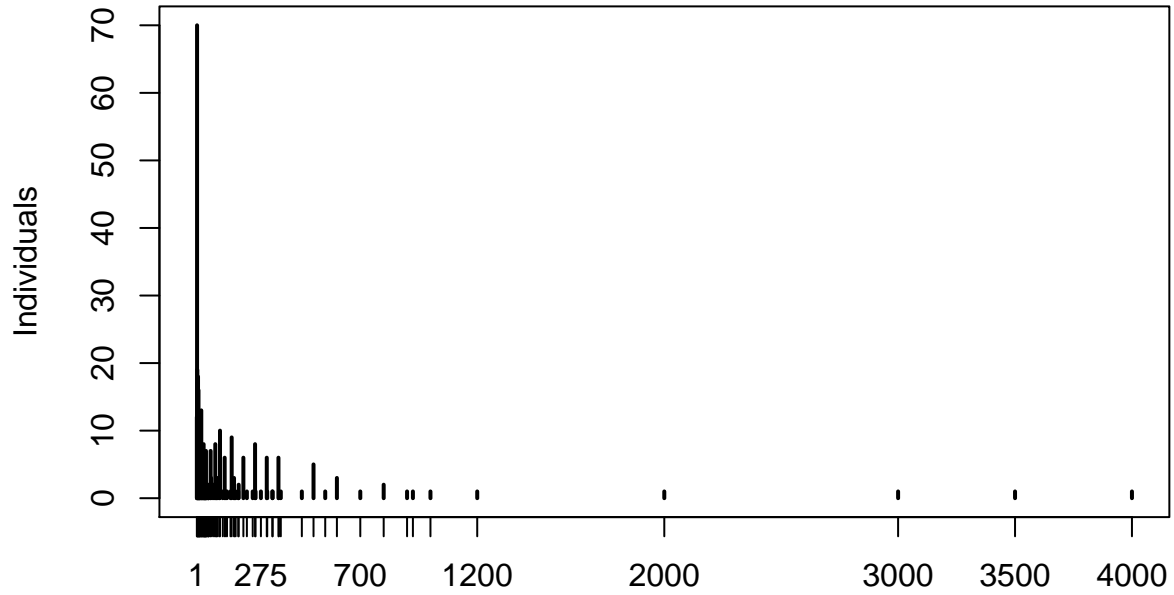


Figure 4: *Enicognathus leptorhynchus* count frequencies

### Estimating effective detection radius (EDR)

We follow the same procedure used for the austral parakeet above.

Table 17: *E. leptorhynchus* EDR models AIC

	df	AIC	dAIC
EDR.avggrouppsize.habitat	4	1054.464	0.00
EDR.habitat.avggrouppsize.numbergroups	5	1059.062	4.60

	df	AIC	dAIC
EDR.avggrouppsize	2	1063.875	9.41
EDR.avggrouppsize.numbergroups	3	1064.553	10.09
EDR.habitat.numbergroups	4	1069.376	14.91
EDR.habitatype	3	1070.755	16.29
EDR.null	1	1086.330	31.87
EDR.numbergroups	2	1087.172	32.71

The model (**EDR.avggrouppsize.habitat**) has the lowest AIC (Table 16), indicating that habitat type and average group size affect the effective detection radius (EDR) of *Enicognathus leptorhynchus*. The mean EDR for each habitat, predicted using model coefficients and the habitat-specific means of average group size, is shown in Table 17.

Table 18: *E. leptorhynchus* top-ranked EDR model estimates

	Estimate	Std. Error	z value	Pr(> z )
log.tau_(Intercept)	5.361	0	1.273654e+51	0
log.tau_gavg	0.001	0	9.610881e+50	0
log.tau_Urban	-0.266	0	-6.308165e+49	0
log.tau_Agropastoral	-0.029	0	-4.843080e+48	0

Table 19: *E. leptorhynchus* habitat-specific mean EDR (m)

Habitat	EDR
Other	235.4183
Urban	224.2124
Agropastoral	180.7002

## Models for number of groups

The model for number of groups is  $G_i \sim \text{Poisson}(D_i A_i)$ , where  $D_i$  = covariates and  $A_i$  = area sampled in site  $i$ .  $A_i$  is calculated using the habitat-specific estimated EDR, and is added to the model as an offset.

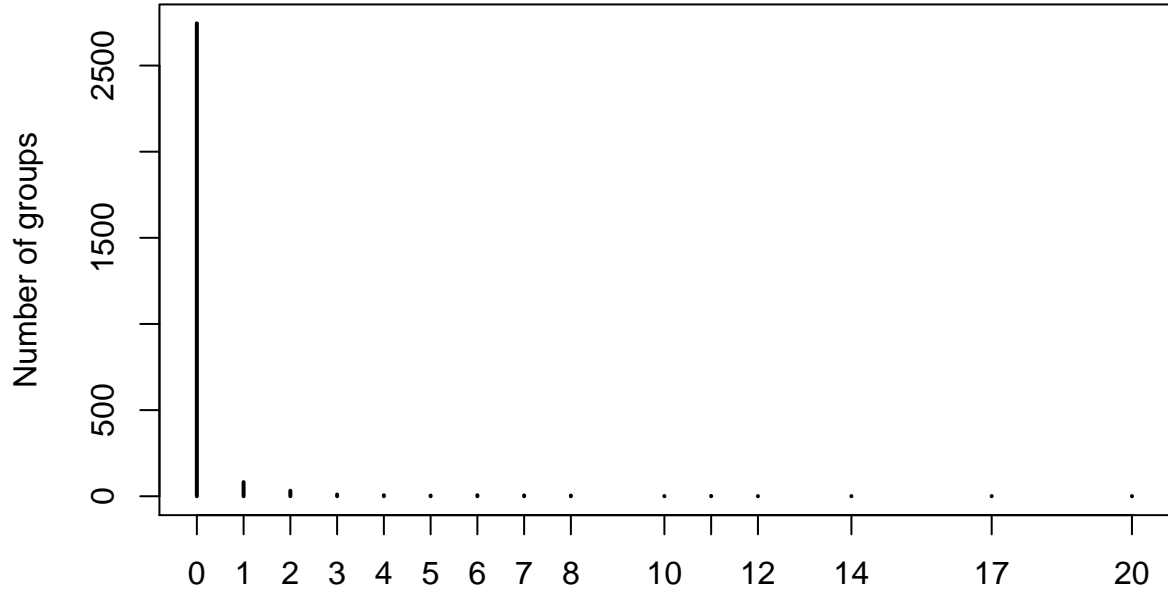


Figure 5: *Enicognathus leptorhynchus* group numbers

### Model selection

First set of models to evaluate the effect of habitat type and elevation covariates (Table 18):

Table 20: *E. leptorhynchus* number of group models

	df	AIC	dAIC
ngroup.hab.ele2	5	2486.990	0.00
ngroup.hab.ele	4	2487.834	0.84
ngroup.hab	3	2493.346	6.36
ngroup.ele2	3	2746.031	259.04
ngroup.ele	2	2761.122	274.13

The models ‘ngroup.hab.ele’ (with linear elevation effect) and ‘ngroup.hab.ele2’ (with quadratic effect) are equally parsimonious (Table 18), but the quadratic term in the latter is mostly uninformative (Tables 19-20). Given these are nested models, we drop the quadratic term and continue with the model with habitat and linear elevation effects.

Table 21: *E. leptorhynchus* ‘ngroup.hab.ele’ model estimates

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-2.781	0.123	-22.568	0.000
habitatOther	-2.035	0.147	-13.874	0.000
habitatUrban	-0.734	0.143	-5.116	0.000
elevation	0.000	0.000	2.736	0.006

Table 22: Deviance partitioning of ‘ngroup.hab.ele’ model for *E. leptorhynchus*

	Df	Deviance	Resid. Df	Resid. Dev
NULL	NA	NA	2900	2358.856
habitat	2	270.326	2898	2088.530
elevation	1	7.512	2897	2081.018

Adding within-year temporal covariates (breeding/non-breeding season and Julian date) and their interactions with habitat:

Table 23: *E. leptorhynchus* number of group models (within-year temporal predictors) AIC

	df	AIC	dAIC
ngroup.habXjdate.ele	7	2110.653	0.00
ngroup.hab.ele.jdate	5	2230.345	119.69
ngroup.habXseason.ele	7	2472.731	362.08
ngroup.hab.ele	4	2487.834	377.18
ngroup.hab.ele.season	5	2489.059	378.41

The model ‘ngroup.habXjdate.ele’ has the lowest AIC (Table 21).

Table 24: *E. leptorhynchus* ‘ngroup.habXjdate.ele’ model estimates

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-1.463	0.167	-8.741	0.000
elevation	0.000	0.000	1.939	0.053
habitatOther	2.234	0.590	3.788	0.000
habitatUrban	0.412	0.310	1.329	0.184
jdate	-0.006	0.001	-10.029	0.000
habitatOther:jdate	-0.040	0.007	-5.589	0.000
habitatUrban:jdate	-0.009	0.002	-4.095	0.000

Finally, we assess year effects by adding a year covariate (2013-2016).

Table 25: *E. leptorhynchus* number of group models (year predictor)  
AIC table

	df	AIC	dAIC
ngroup.habXjdate.ele.year	10	1907.516	0.00
ngroup.habXjdate.ele	7	2110.653	203.14

The best-ranked model indicates that the number of groups is affected by habitat type, elevation, Julian date and year (Tables 23-24).

Table 26: *E. leptorhynchus* ‘ngroup.habXjdate.ele.year’ model estimates

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-3.424	0.553	-6.189	0.000
elevation	0.001	0.000	3.208	0.001
habitatOther	0.677	0.350	1.932	0.053
habitatUrban	0.566	0.257	2.206	0.027
jdate	-0.002	0.002	-1.291	0.197
as.factor(year)2014	-0.201	0.352	-0.570	0.569
as.factor(year)2015	1.388	0.233	5.953	0.000
as.factor(year)2016	1.909	0.442	4.322	0.000
habitatOther:jdate	-0.020	0.004	-5.091	0.000
habitatUrban:jdate	-0.008	0.002	-4.327	0.000

Table 27: Deviance partitioning of ‘ngroup.habXjdate.ele.year’ model for *E. leptorhynchus*

	Df	Deviance	Resid. Df	Resid. Dev
NULL	NA	NA	2900	2358.856
elevation	1	0.551	2899	2358.306
habitat	2	277.287	2897	2081.018
jdate	1	259.490	2896	1821.528
as.factor(year)	3	239.982	2893	1581.547
habitat:jdate	2	92.847	2891	1488.700

## Models for group size

### Goodness-of-fit

Model goodness-of-fit can be evaluated visually, inspecting the proportions of fitted values against the count distribution. Deviation from the 1:1 line can be used as a goodness-of-fit metric, with the better model showing smaller deviation.




Figure 6: *Enicognathus leptorhynchus* group size model goodness-of-fit (red: null model; blue: habitat.season model). Better model is closer to 1:1 line.

The deviation of null model is 1.62, and of the VIP.habitat.season model is 1.65.

#### Confidence intervals for model estimates

Table 28: *E. leptorhynchus* group size model estimates

	Estimate	Std. Error	z value	Pr(> z )
P_(Intercept)	-8.066	NA	NA	NA
P_Urban	8.809	NA	NA	NA
P_Agropastoral	31.024	NA	NA	NA
V_(Intercept)	22.401	NA	NA	NA
V_seasonnon-breeding	-17.695	NA	NA	NA

Model-fitting function cannot estimate coefficient standard errors due to singular Hessian matrix (Table 25). We can calculate confidence intervals based on quantiles using bootstrap (with n=250):

Table 29: Estimated coefficients and 95% CI based on bootstrap quantiles for *E. leptorhynchus* group size models

	Estimate	2.5%	97.5%
P_(Intercept)	-8.066	-20.905	30.686
P_Urban	8.809	-26.727	26.113
P_Agropastoral	31.024	-25.722	26.283
V_(Intercept)	22.401	-24.223	25.795
V_seasonnon-breeding	-17.695	-25.549	27.828

Results indicate that habitat type and season are not good predictors of group size for the slender-billed parakeet.

## Density Predictions

The mean expected abundance ( $E$ ) can be predicted for each surveyed site  $i$  as  $E_i = (\phi_i 2 + (1 - \phi_i) \mu_i) N_i$ , where  $N$  is the mean number of groups,  $\mu$  is the mean group size and  $\phi$  is the V(=2)-inflation probability (i.e. the probability that a given group will be a pair).

We can then obtain the expected density stratified by habitat type, year and season, as  $D = \frac{\sum E}{\sum A}$ , where  $A$  is the area surveyed.

Because habitat and season are not good predictors of group size for the Slender-billed Parakeet, we predict density using the  $\mu$  and  $\phi$  estimates from the null model for Gsize.

Table 30: *E. ferrugineous* mean density predictions for each habitat type, year and season

Habitat	Year	Season	Density (individuals/ha)	Density (2.5% CI)	Density (97.5% CI)
Agropastoral	2013	breeding	0.059	0.038	0.092
Other	2013	breeding	0.000	0.000	0.002
Urban	2013	breeding	0.010	0.004	0.022
Agropastoral	2015	breeding	0.258	0.201	0.326
Other	2015	breeding	0.001	0.000	0.008
Urban	2015	breeding	0.036	0.017	0.079
Agropastoral	2013	non-breeding	0.073	0.048	0.112
Other	2013	non-breeding	0.000	0.000	0.003
Urban	2013	non-breeding	0.013	0.006	0.028
Agropastoral	2014	non-breeding	0.075	0.052	0.110
Other	2014	non-breeding	0.014	0.008	0.025
Urban	2014	non-breeding	0.044	0.028	0.068
Agropastoral	2015	non-breeding	0.295	0.234	0.369
Other	2015	non-breeding	0.003	0.001	0.015
Urban	2015	non-breeding	0.058	0.032	0.110
Agropastoral	2016	non-breeding	0.795	0.652	0.973
Other	2016	non-breeding	0.486	0.364	0.661
Urban	2016	non-breeding	0.730	0.531	0.990

# VIP model - R functions and simulations

Peter Solymos

## Functions

The `vip` function does the optimization. `method` argument can take values listed for `optim` and also "DE" for differential evolution algorithm. If there are convergence issues with "Nelder-Mead", try "SANN" and "DE".

```
library(DEoptim)
library(Matrix)
vip <-
function(Y, X, Z, V=0,
offsetx, offsetz, weights, linkz="logit",
truncate=FALSE, hessian=TRUE, method="Nelder-Mead", init=NULL, ...) {
  if (missing(Y))
    stop("C'mon, you must have some data?!")
  if (truncate && any(Y < 1))
    stop("Y must be >0 when truncate=TRUE")
  n <- length(Y)
  id0 <- Y == V
  id1 <- !id0
  if (missing(X)) {
    X <- matrix(1, n, 1)
    colnames(X) <- "(Intercept)"
  }
  if (missing(Z)) {
    Z <- matrix(1, n, 1)
    colnames(Z) <- "(Intercept)"
  }
  kx <- ncol(X)
  kz <- ncol(Z)
  if (missing(offsetx))
    offsetx <- 0
  if (missing(offsetz))
    offsetz <- 0
  if (missing(weights))
    weights <- rep(1, n)
  linkinvx <- poisson("log")$linkinv
  linkinvz <- binomial(linkz)$linkinv
  good.num.limit <- c(.Machine$double.xmin, .Machine$double.xmax)^(1/3)

  ## VIP model full likelihood
  nll_VIP_ML <- function(parms) {
    mu <- as.vector(linkinvx(X %*% parms[1:kx] + offsetx))
    phi <- as.vector(linkinvz(Z %*% parms[(kx + 1):(kx + kz)] + offsetz))
    loglik0 <- log(phi + (1 - phi) * dpois(V, lambda = mu, log = FALSE))
    loglik1 <- log(1 - phi) + dpois(Y, lambda = mu, log = TRUE)
    loglik <- sum(weights[id0] * loglik0[id0]) + sum(weights[id1] * loglik1[id1])
    if (!is.finite(loglik) || is.na(loglik))
      loglik <- -good.num.limit[2]
    -loglik
  }
  ## 0-truncated VIP model full likelihood
```



```

nll_VIP_TR <- function(parms) {
  mu <- as.vector(linkinvx(X %*% parms[1:kx] + offsetx))
  phi <- as.vector(linkinvz(Z %*% parms[(kx + 1):(kx + kz)] + offsetz))
  loglik0 <- log(phi + (1 - phi) * dpois(V, lambda = mu, log = FALSE) / (1-exp(-mu)))
  loglik1 <- log((1 - phi) * dpois(Y, lambda = mu, log = FALSE) / (1-exp(-mu)))
  loglik <- sum(weights[id0] * loglik0[id0]) + sum(weights[id1] * loglik1[id1])
  if (!is.finite(loglik) || is.na(loglik))
    loglik <- -good.num.limit[2]
  -loglik
}

.solvencnear <-
function(x)
{
  xinv <- try(solve(x), silent = TRUE)
  if (inherits(xinv, "try-error"))
    xinv <- as.matrix(solve(Matrix::nearPD(x)$mat))
  xinv
}

if (is.null(init))

  init <- rep(0, kx+kz)
nll <- if (truncate) nll_VIP_TR else nll_VIP_ML

if (method == "DE") {
  DELimit <- 10
  up <- rep(DELimit, length(init))
  lo <- -up
  opt <- DEoptim(fn=nll, lower=lo, upper=up,
    control=list(trace=FALSE, itermax=length(init)*200))
  par <- opt$optim$bestmem
  names(par) <- c(paste0("P_", colnames(X)), paste0("V_", colnames(Z)))
  ll <- -opt$optim$bestval
  if (hessian) {
    hess <- optimHess(opt$optim$bestmem, nll)
    vc <- .solvencnear(hess)
  } else {
    matrix(NA, length(par), length(par))
  }
} else {
  opt <- optim(init, nll,
    hessian=hessian, method=method, ...)
  par <- opt$par
  names(par) <- c(paste0("P_", colnames(X)), paste0("V_", colnames(Z)))
  vc <- if (hessian)
    .solvencnear(opt$hessian) else matrix(NA, length(par), length(par))
  ll <- -opt$value
}

dimnames(vc) <- list(names(par), names(par))
out <- list(call=match.call(),
  coefficients=par, loglik=ll, vcov=vc, nobs=n,
  truncate=truncate, Y=Y, X=X, Z=Z, V=V,
  offsetx=offsetx, offsetz=offsetz, weights=weights,
  linkz=linkz, method=method, init=init)

```

```

class(out) <- "vip"
out
}
vcov.vip <- function(object, ...) object$vcov
logLik.vip <- function(object, ...)
  structure(object$loglik, df = object$nobs - length(object$coef),
    nobs = object$nobs, class = "logLik")
summary.vip <- function(object, ...) {
  k <- length(object$coefficients)
  coefs <- coef(object)
  se <- sqrt(diag(vcov(object)))
  tstat <- coefs/se
  pval <- 2 * pnorm(-abs(tstat))
  coefs <- cbind(coefs, se, tstat, pval)
  colnames(coefs) <- c("Estimate", "Std. Error", "z value", "Pr(>|z|)")
  coefs <- coefs[1:k, , drop = FALSE]
  rownames(coefs) <- names(coef(object))
  out <- list(call = object$call, coefficients=coefs, loglik = object$loglik,
    bic=BIC(object), truncate=object$truncate)
  class(out) <- "summary.vip"
  return(out)
}
print.summary.vip <- function(x, digits, ...)
{
  if (missing(digits))
    digits <- max(3, getOption("digits") - 3)
  cat("\nCall:", deparse(x$call,
    width.cutoff = floor(getOption("width") * 0.85)), "", sep = "\n")
  cat("V-Inflated", if (x$truncate) "(Zero-Truncated)" else "", "Poisson Model\n\n")
  cat(paste("Coefficients:\n", sep = ""))
  printCoefmat(x$coefficients, digits = digits, signif.legend = FALSE)
  if (!any(is.na(array(x$coefficients)))) {
    if (getOption("show.signif.stars") & any(x$coefficients[,4] < 0.1))
      cat("----\nSignif. codes: ", "0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1", "\n")
  }
  cat("\nLog-likelihood:", formatC(x$loglik, digits = digits),
    "\nBIC =", formatC(x$bic, digits = digits), "\n")
  cat("\n")
  invisible(x)
}
confint.vip <-
function(object, parm, level = 0.95, ...)
{
  cf <- coef(object)
  pnames <- names(cf)
  if (missing(parm)) {
    parm <- pnames
  } else {
    if (is.numeric(parm))
      parm <- pnames[parm]
  }
  a <- (1 - level)/2
  a <- c(a, 1 - a)

```

```

pct <- paste(format(100 * a, trim = TRUE, scientific = FALSE, digits = 3), "%", sep="")
ci <- array(NA, dim = c(length(parm), 2), dimnames = list(parm, pct))
fac <- qnorm(a)
ses <- sqrt(diag(vcov(object, model, type)))
ci[] <- cf[parm] + ses[parm] %o% fac
ci
}

```

## Simple case

```

set.seed(123)
n <- 1000
lam <- 2 # poisson mean, can be a vector of length n
phi <- 0.4 # V-inflation probability, can be a vector of length n
V <- 2 # V is the count value, can be 0, 2, etc
y <- y0 <- rpois(n, lam)
a <- rbinom(n, 1, phi)
y[a > 0] <- V
table(Poisson=y0, Vinflated=y)

```

```

##          Vinflated
## Poisson  0  1  2  3  4  5  6  8
##          0 81  0 51  0  0  0  0
##          1  0 151 126  0  0  0  0
##          2  0  0 274  0  0  0  0
##          3  0  0  65 112  0  0  0
##          4  0  0  39  0 43  0  0
##          5  0  0  12  0  0 29  0
##          6  0  0  6  0  0  0  9
##          7  0  0  1  0  0  0  0
##          8  0  0  0  0  0  0  1

```

```

mod <- vip(Y=y, V=2)
summary(mod)

```

```

##
## Call:
## vip(Y = y, V = 2)
##
## V-Inflated Poisson Model
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## P_(Intercept)  0.70472    0.02909  24.224 < 2e-16 ***
## V_(Intercept) -0.33900    0.08824  -3.842 0.000122 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Log-likelihood: -1345
## BIC = 9585

```

```

cbind(True=c(log_lam=log(lam), logit_phi=qlogis(phi)),
      Est=coef(mod))

```

```
##               True      Est
## log_lam      0.6931472  0.7047243
## logit_phi    -0.4054651 -0.3389963
```

## Covariates for the non-V part

```
set.seed(123)
n <- 1000
x <- rnorm(n)
df <- data.frame(x=x)
X <- model.matrix(~x, df)
beta <- c(-0.5,-0.5) # Intercept and beta values for covariate
lam <- exp(X %*% beta) # poisson mean, can be a vector of length n
phi <- 0.4 # V-inflation probability, can be a vector of length n
V <- 2 # V is the count value, can be 0, 2, etc
y <- y0 <- rpois(n, lam)
a <- rbinom(n, 1, phi)
y[a > 0] <- V
table(Poisson=y0, Vinflated=y)
```

```
##           Vinflated
## Poisson   0   1   2   3   4   5   6
##           0 299   0 228   0   0   0   0
##           1   0 201 131   0   0   0   0
##           2   0   0  95   0   0   0   0
##           3   0   0  14  18   0   0   0
##           4   0   0   2   0   6   0   0
##           5   0   0   0   0   0   4   0
##           6   0   0   0   0   0   0   2
```

```
mod <- vip(Y=y, X=X, V=2)
summary(mod)
```

```
##
## Call:
## vip(Y = y, X = X, V = 2)
##
## V-Inflated Poisson Model
##
## Coefficients:
##           Estimate Std. Error z value Pr(>|z|)
## P_(Intercept) -0.39671    0.06382  -6.216 5.11e-10 ***
## P_x           -0.49031    0.05331  -9.198 < 2e-16 ***
## V_(Intercept) -0.42040    0.07813  -5.381 7.42e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Log-likelihood: -1141
## BIC = 9169
```

```
cbind(True=c(beta=beta, logit_phi=qlogis(phi)),
      Est=coef(mod))
```

```
##               True      Est
```

```
## beta1      -0.5000000 -0.3967122
## beta2      -0.5000000 -0.4903124
## logit_phi  -0.4054651 -0.4203965
```

## Methods

```
coef(mod)
```

```
## P_(Intercept)      P_x V_(Intercept)
##      -0.3967122      -0.4903124      -0.4203965
```

```
vcov(mod)
```

```
##              P_(Intercept)      P_x V_(Intercept)
## P_(Intercept)  0.004073534  0.0018566813 -0.0014632502
## P_x           0.001856681  0.0028417922 -0.0006105607
## V_(Intercept) -0.001463250 -0.0006105607  0.0061042545
```

```
summary(mod)
```

```
##
## Call:
## vip(Y = y, X = X, V = 2)
##
## V-Inflated Poisson Model
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## P_(Intercept) -0.39671      0.06382  -6.216 5.11e-10 ***
## P_x           -0.49031      0.05331  -9.198 < 2e-16 ***
## V_(Intercept) -0.42040      0.07813  -5.381 7.42e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Log-likelihood: -1141
## BIC = 9169
```

```
confint(mod)
```

```
##              2.5%      97.5%
## P_(Intercept) -0.5218054 -0.2716189
## P_x           -0.5947950 -0.3858297
## V_(Intercept) -0.5735279 -0.2672650
```

```
nobs(mod)
```

```
## [1] 1000
```

```
logLik(mod)
```

```
## 'log Lik.' -1140.751 (df=997)
```

```
AIC(mod)
```

```
## [1] 4275.502
```

```
BIC(mod)
```

```
## [1] 9168.534
```

## Zero-truncated VIP

We can truncate counts to be larger than 0. We also need  $V > 0$  (for  $V = 0$  case, look into ZIP or conditional Poisson model). Conceptually, the V-Inflation follows the 0-truncation (because we cannot observe 0, real truncated distribution).

The 0-truncated PDF is  $P(Y = y \mid Y > 0) = \frac{P(Y=y)}{1-P(Y=0)}$ . The 0-truncated V-Inflated density is  $P(Y = y \mid Y > 0, V > 0) = \phi I(Y = V) + (1 - \phi) \frac{f(y;\lambda)}{1-f(0;\lambda)}$ . This can be achieved in the `vip` call by the argument `truncate=TRUE`.

Here we use covariates for both the V and non-V part.

```
set.seed(1)
n <- 1000
x <- rnorm(n)
z <- runif(n, -1, 1)
df <- data.frame(x=x, z=z)
X <- model.matrix(~x, df)
Z <- model.matrix(~z, df)
beta <- c(-0.5, -0.5)
alpha <- c(0, 0.5)
lam <- exp(X %*% beta)
phi <- plogis(Z %*% alpha)
V <- 2 # V is the count value, cannot be 0
y <- y0 <- rpois(n, lam)
a <- rbinom(n, 1, phi)
keep <- y0>0
y <- y[keep] # conditioning (i.e. exclude 0s)
y0 <- y0[keep]
X <- X[keep,]
Z <- Z[keep,]
y[a[keep] > 0] <- V
table(Poisson=y0, Vinflated=y)
```

```
##          Vinflated
## Poisson   1   2   3   4   6
##         1 155 141   0   0   0
##         2   0 127   0   0   0
##         3   0  21  16   0   0
##         4   0   4   0   7   0
##         5   0   2   0   0   0
##         6   0   0   0   0   1
```

```
mod <- vip(Y=y, X=X, Z=Z, V=2, truncate=TRUE)
summary(mod)
```

```
##
## Call:
## vip(Y = y, X = X, Z = Z, V = 2, truncate = TRUE)
##
## V-Inflated (Zero-Truncated) Poisson Model
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## P_(Intercept) -0.50814    0.14803  -3.433 0.000598 ***
## P_x           -0.57344    0.10170  -5.639 1.71e-08 ***
```

```
## V_(Intercept) 0.02131 0.12572 0.170 0.865387
## V_z          0.47041 0.20691 2.273 0.022999 *
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Log-likelihood: -384.1
## BIC = 3664

cbind(True=c(beta=beta, alpha=alpha),
      Est=coef(mod))

##           True           Est
## beta1  -0.5 -0.50813933
## beta2  -0.5 -0.57343540
## alpha1  0.0  0.02131236
## alpha2  0.5  0.47040670
```

## Goodness of fit

```
goodness <- function(object, maxcount=NULL) {
  if (is.null(maxcount))
    maxcount <- max(object$Y)
  COUNTS <- if (object$truncate)
    1L:max(object$Y) else 0L:maxcount

  P_obs <- table(object$Y) / nobs(object)
  P_obs <- as.numeric(P_obs[match(COUNTS, names(P_obs))])
  P_obs[is.na(P_obs)] <- 0
  names(P_obs) <- COUNTS

  P_exp <- P_obs
  P_exp[] <- 0

  linkinvx <- poisson("log")$linkinv
  linkinvz <- binomial(object$linkz)$linkinv
  parms <- coef(object)
  kx <- ncol(object$X)
  kz <- ncol(object$Z)
  mu <- as.vector(linkinvx(object$X %*% parms[1:kx] + object$offsetx))
  phi <- as.vector(linkinvz(object$Z %*% parms[(kx + 1):(kx + kz)] + object$offsetz))
  #id0 <- object$Y == object$V
  Pmat <- matrix(0, nobs(object), length(COUNTS))
  colnames(Pmat) <- COUNTS

  PV <- if (object$truncate) {
    log(phi + (1 - phi) * dpois(object$V,
      lambda = mu, log = FALSE) / (1-exp(-mu)))
  } else {
    log(phi + (1 - phi) * dpois(object$V, lambda = mu, log = FALSE))
  }
  for (i in COUNTS) {
    PC <- if (object$truncate) {
      log((1 - phi) * dpois(i, lambda = mu, log = FALSE) / (1-exp(-mu)))
    }
```

```

    } else {
      log(1 - phi) + dpois(i, lambda = mu, log = TRUE)
    }
    Pmat[,as.character(i)] <- if (i == object$V)
      exp(PV) else exp(PC)
  }
  P_exp <- colMeans(Pmat)
  cbind(P_obs=P_obs, P_exp=P_exp)
}

## fit null model
mod0 <- vip(Y=y, X=X[,1,drop=FALSE], Z=Z[,1,drop=FALSE], V=2, truncate=TRUE)
## calculate GoF for null and other model
(gof0 <- goodness(mod0))

##          P_obs          P_exp
## 1 0.327004219 0.3234032597
## 2 0.622362869 0.6223385125
## 3 0.033755274 0.0427550867
## 4 0.014767932 0.0095197408
## 5 0.000000000 0.0016957134
## 6 0.002109705 0.0002517089

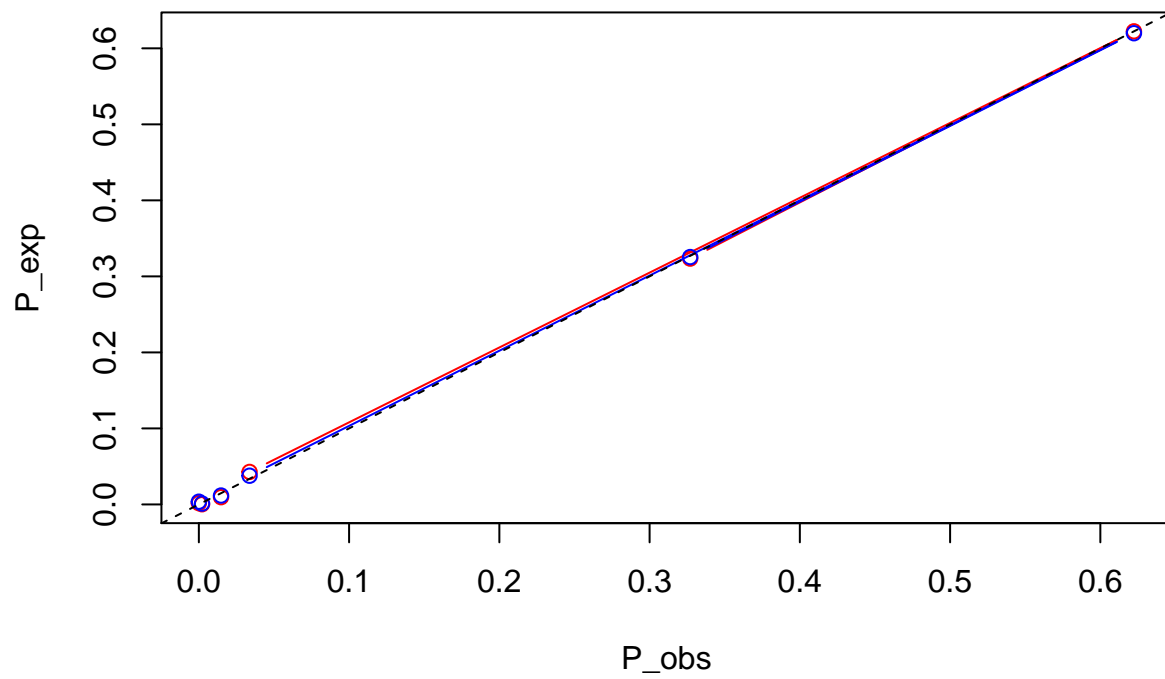
(gof <- goodness(mod))

##          P_obs          P_exp
## 1 0.327004219 0.325425662
## 2 0.622362869 0.619861184
## 3 0.033755274 0.037772839
## 4 0.014767932 0.011711772
## 5 0.000000000 0.003622341
## 6 0.002109705 0.001119974

## better model is closer to the 1:1 line
plot(gof0, type="b", col=2)
points(gof, type="b", col=4)
abline(0, 1, lty=2)

```





```
## better model will give smaller absolute deviation  
sum(abs(apply(gof0, 1, diff)))
```

```
## [1] 0.02142703
```

```
sum(abs(apply(gof, 1, diff)))
```

```
## [1] 0.01576604
```