

V-inflated Poisson count regression

Motivation

We want to model the distribution of groups size in parrot populations. These animals tend to form smaller or larger groups, but groups of size 2 are also often observed as a result of mating pairs.

The count distribution is characterized with a spike at 2, and by the absence of 0s due to group size being conditional on having >0 birds to consider it a group.

We start developing a general V-Inflated Poisson (VIP) model, then we add the >0 condition. Simulations are done to check the estimating procedure.

Maximum likelihood

Let Y be a random variable, and y are observations, V is the count value that has some extra probability mass ($V = 0$ is the ZIP model), $f(y; \lambda)$ is the Poisson density ($f(y; \lambda) = e^{-\lambda} \frac{\lambda^y}{y!}$).

The V-Inflated density can be written as $P(Y = y) = \phi I(Y = V) + (1 - \phi)f(y; \lambda)$ which is $\phi + (1 - \phi)f(V; \lambda)$ when $Y = V$ and $(1 - \phi)f(y; \lambda)$ otherwise.

Functions

```
vip <-  
function(Y, X, Z, V=0,  
offsetx, offsetz, weights, linkz="logit",  
truncate=FALSE, hessian=TRUE, method="Nelder-Mead", ...) {  
  if (missing(Y))  
    stop("C'mon, you must have some data?!")  
  if (truncate && any(Y < 1))  
    stop("Y must be >0 when truncate=TRUE")  
  n <- length(Y)  
  id0 <- Y == V  
  id1 <- !id0  
  if (missing(X)) {  
    X <- matrix(1, n, 1)  
    colnames(X) <- "(Intercept)"  
  }  
  if (missing(Z)) {  
    Z <- matrix(1, n, 1)  
    colnames(Z) <- "(Intercept)"  
  }  
  kx <- ncol(X)  
  kz <- ncol(Z)  
  if (missing(offsetx))  
    offsetx <- 0  
  if (missing(offsetz))  
    offsetz <- 0  
  if (missing(weights))  
    weights <- rep(1, n)
```

```

linkinvx <- poisson("log")$linkinv
linkinvz <- binomial(linkz)$linkinv
good.num.limit <- c(.Machine$double.xmin, .Machine$double.xmax)^(1/3)

## VIP model full likelihood
nll_VIP_ML <- function(parms) {
  mu <- as.vector(linkinvx(X %*% parms[1:kx] + offsetx))
  phi <- as.vector(linkinvz(Z %*% parms[(kx + 1):(kx + kz)] + offsetz))
  loglik0 <- log(phi + (1 - phi) * dpois(V, lambda = mu, log = FALSE))
  loglik1 <- log(1 - phi) + dpois(Y, lambda = mu, log = TRUE)
  loglik <- sum(weights[id0] * loglik0[id0]) + sum(weights[id1] * loglik1[id1])
  if (!is.finite(loglik) || is.na(loglik))
    loglik <- -good.num.limit[2]
  -loglik
}

## 0-truncated VIP model full likelihood
nll_VIP_TR <- function(parms) {
  mu <- as.vector(linkinvx(X %*% parms[1:kx] + offsetx))
  phi <- as.vector(linkinvz(Z %*% parms[(kx + 1):(kx + kz)] + offsetz))
  loglik0 <- log(phi + (1 - phi) * dpois(V, lambda = mu, log = FALSE) / (1-exp(-mu)))
  loglik1 <- log((1 - phi) * dpois(Y, lambda = mu, log = FALSE) / (1-exp(-mu)))
  loglik <- sum(weights[id0] * loglik0[id0]) + sum(weights[id1] * loglik1[id1])
  if (!is.finite(loglik) || is.na(loglik))
    loglik <- -good.num.limit[2]
  -loglik
}

opt <- optim(rep(0, kx+kz),
  if (truncate) nll_VIP_TR else nll_VIP_ML,
  hessian=hessian, method=method, ...)
par <- opt$par
names(par) <- c(paste0("P_", colnames(X)), paste0("V_", colnames(Z)))
vc <- if (hessian)
  solve(opt$hessian) else matrix(NA, length(par), length(par))
dimnames(vc) <- list(names(par), names(par))
out <- list(call=match.call(),
  coefficients=par, loglik=-opt$value, vcov=vc, nobs=n,
  truncate=truncate)
class(out) <- "vip"
out
}

vcov.vip <- function(object, ...) object$vcov
logLik.vip <- function(object, ...)
  structure(object$loglik, df = object$nobs - length(object$coef),
    nobs = object$nobs, class = "logLik")
summary.vip <- function(object, ...) {
  k <- length(object$coefficients)
  coefs <- coef(object)
  se <- sqrt(diag(vcov(object)))
  tstat <- coefs/se
  pval <- 2 * pnorm(-abs(tstat))
  coefs <- cbind(coefs, se, tstat, pval)
  colnames(coefs) <- c("Estimate", "Std. Error", "z value", "Pr(>|z|)")
}

```

```

  coefs <- coefs[1:k, , drop = FALSE]
  rownames(coefs) <- names(coef(object))
  out <- list(call = object$call, coefficients=coefs, loglik = object$loglik,
    bic=BIC(object), truncate=object$truncate)
  class(out) <- "summary.vip"
  return(out)
}

print.summary.vip <- function (x, digits, ...)
{
  if (missing(digits))
    digits <- max(3, getOption("digits") - 3)
  cat("\nCall:", deparse(x$call,
    width.cutoff = floor(getOption("width") * 0.85)), "", sep = "\n")
  cat("\nV-Inflated", if (x$truncate) "(Zero-Truncated)" else "", "Poisson Model\n\n")
  cat(paste("Coefficients:\n", sep = ""))
  printCoefmat(x$coefficients, digits = digits, signif.legend = FALSE)
  if (!any(is.na(array(x$coefficients)))) {
    if (getOption("show.signif.stars") & any(x$coefficients[,4] < 0.1))
      cat("----\nSignif. codes: ", "0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1", "\n")
  }
  cat("\nLog-likelihood:", formatC(x$loglik, digits = digits),
    "\nBIC =", formatC(x$bic, digits = digits), "\n")
  cat("\n")
  invisible(x)
}

confint.vip <-
function (object, parm, level = 0.95, ...)
{
  cf <- coef(object)
  pnames <- names(cf)
  if (missing(parm)) {
    parm <- pnames
  } else {
    if (is.numeric(parm))
      parm <- pnames[parm]
  }
  a <- (1 - level)/2
  a <- c(a, 1 - a)
  pct <- paste(format(100 * a, trim = TRUE, scientific = FALSE, digits = 3), "%", sep="")
  ci <- array(NA, dim = c(length(parm), 2), dimnames = list(parm, pct))
  fac <- qnorm(a)
  ses <- sqrt(diag(vcov(object, model, type)))
  ci[] <- cf[parm] + ses[parm] %o% fac
  ci
}

```

Simple case

```

set.seed(123)
n <- 1000
lam <- 2 # poisson mean, can be a vector of length n
phi <- 0.4 # V-inflation probability, can be a vector of length n

```

```
V <- 2 # V is the count value, can be 0, 2, etc
y <- y0 <- rpois(n, lam)
a <- rbinom(n, 1, phi)
y[a > 0] <- V
table(Poisson=y0, Vinflated=y)
```

```
##          Vinflated
## Poisson  0  1  2  3  4  5  6  8
##      0 81  0 51  0  0  0  0  0
##      1  0 151 126  0  0  0  0  0
##      2  0  0 274  0  0  0  0  0
##      3  0  0  65 112  0  0  0  0
##      4  0  0  39  0 43  0  0  0
##      5  0  0  12  0  0 29  0  0
##      6  0  0  6  0  0  0  9  0
##      7  0  0  1  0  0  0  0  0
##      8  0  0  0  0  0  0  0  1
```

```
mod <- vip(Y=y, V=2)
summary(mod)
```

```
##
## Call:
## vip(Y = y, V = 2)
##
## V-Inflated Poisson Model
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## P_(Intercept)  0.70472    0.02909  24.224 < 2e-16 ***
## V_(Intercept) -0.33900    0.08824  -3.842 0.000122 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Log-likelihood: -1345
## BIC = 9585
```

```
cbind(True=c(log_lam=log(lam), logit_phi=qlogis(phi)),
      Est=coef(mod))
```

```
##              True      Est
## log_lam    0.6931472 0.7047243
## logit_phi -0.4054651 -0.3389963
```

Covariates for the non-V part

```
set.seed(123)
n <- 10000
x <- rnorm(n)
df <- data.frame(x=x)
X <- model.matrix(~x, df)
beta <- c(-0.5, -0.5) # Intercept and beta values for covariate
lam <- exp(X %*% beta) # poisson mean, can be a vector of length n
phi <- 0.4 # V-inflation probability, can be a vector of length n
```

```
V <- 2 # V is the count value, can be 0, 2, etc
y <- y0 <- rpois(n, lam)
a <- rbinom(n, 1, phi)
y[a > 0] <- V
table(Poisson=y0, Vinflated=y)
```

```
##          Vinflated
## Poisson    0     1     2     3     4     5     6     7     8
##      0 3182     0 2131     0     0     0     0     0     0
##      1     0 1981 1137     0     0     0     0     0     0
##      2     0     0 1088     0     0     0     0     0     0
##      3     0     0  118  226     0     0     0     0     0
##      4     0     0   40     0    57     0     0     0     0
##      5     0     0   14     0     0    17     0     0     0
##      6     0     0    1     0     0     0     3     0     0
##      7     0     0    2     0     0     0     0     1     0
##      8     0     0    1     0     0     0     0     0     1
```

```
mod <- vip(Y=y, X=X, V=2)
summary(mod)
```

```
##
## Call:
## vip(Y = y, X = X, V = 2)
##
## V-Inflated Poisson Model
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## P_(Intercept) -0.45313      0.02037  -22.24  <2e-16 ***
## P_x           -0.49231      0.01664  -29.58  <2e-16 ***
## V_(Intercept) -0.48770      0.02483  -19.64  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Log-likelihood: -1.133e+04
## BIC = 1.147e+05
```

```
cbind(True=c(beta=beta, logit_phi=qlogis(phi)),
      Est=coef(mod))
```

```
##              True      Est
## beta1      -0.5000000 -0.4531273
## beta2      -0.5000000 -0.4923100
## logit_phi  -0.4054651 -0.4876957
```

Methods

```
coef(mod)
```

```
## P_(Intercept)          P_x V_(Intercept)
##      -0.4531273      -0.4923100      -0.4876957
```

```
vcov(mod)
```

```
##               P_(Intercept)          P_x V_(Intercept)
## P_(Intercept) 0.0004151059 1.815322e-04 -1.454395e-04
## P_x           0.0001815322 2.769780e-04 -5.339019e-05
## V_(Intercept) -0.0001454395 -5.339019e-05 6.165031e-04

summary(mod)

##
## Call:
## vip(Y = y, X = X, V = 2)
##
## V-Inflated Poisson Model
##
## Coefficients:
##               Estimate Std. Error z value Pr(>|z|)
## P_(Intercept) -0.45313    0.02037  -22.24  <2e-16 ***
## P_x           -0.49231    0.01664  -29.58  <2e-16 ***
## V_(Intercept) -0.48770    0.02483  -19.64  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Log-likelihood: -1.133e+04
## BIC = 1.147e+05

confint(mod)

##               2.5%          97.5%
## P_(Intercept) -0.4930599 -0.4131947
## P_x           -0.5249290 -0.4596910
## V_(Intercept) -0.5363606 -0.4390308

nobs(mod)

## [1] 10000

logLik(mod)

## 'log Lik.' -11332.89 (df=9997)

AIC(mod)

## [1] 42659.77

BIC(mod)

## [1] 114741.5
```

Zero-truncated VIP

We can truncate counts to be larger than 0. We also need $V > 0$ (for $V = 0$ case, look into ZIP or conditional Poisson model). Conceptually, the V-Inflation follows the 0-truncation (because we cannot observe 0, real truncated distribution).

The 0-truncated PDF is $P(Y = y \mid Y > 0) = \frac{P(Y=y)}{1-P(Y=0)}$. The 0-truncated V-Inflated density is $P(Y = y \mid Y > 0, V > 0) = \phi I(Y = V) + (1 - \phi) \frac{f(y;\lambda)}{1-f(0;\lambda)}$. This can be achieved in the `vip` call by the argument `truncate=TRUE`.

Here we use covariates for both the V and non-V part.

```

set.seed(1)
n <- 1000
x <- rnorm(n)
z <- runif(n, -1, 1)
df <- data.frame(x=x, z=z)
X <- model.matrix(~x, df)
Z <- model.matrix(~z, df)
beta <- c(-0.5, -0.5)
alpha <- c(0, 0.5)
lam <- exp(X %*% beta)
phi <- plogis(Z %*% alpha)
V <- 2 # V is the count value, cannot be 0
y <- y0 <- rpois(n, lam)
a <- rbinom(n, 1, phi)
keep <- y0>0
y <- y[keep] # conditioning (i.e. exclude 0s)
y0 <- y0[keep]
X <- X[keep,]
Z <- Z[keep,]
y[a[keep] > 0] <- V
table(Poisson=y0, Vinflated=y)

```

```

##          Vinflated
## Poisson   1   2   3   4   6
##          1 155 141   0   0   0
##          2   0 127   0   0   0
##          3   0  21  16   0   0
##          4   0   4   0   7   0
##          5   0   2   0   0   0
##          6   0   0   0   0   1

```

```

mod <- vip(Y=y, X=X, Z=Z, V=2, truncate=TRUE)
summary(mod)

```

```

##
## Call:
## vip(Y = y, X = X, Z = Z, V = 2, truncate = TRUE)
##
## V-Inflated (Zero-Truncated) Poisson Model
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## P_(Intercept) -0.50814    0.14803  -3.433 0.000598 ***
## P_x           -0.57344    0.10170  -5.639 1.71e-08 ***
## V_(Intercept)  0.02131    0.12572   0.170 0.865387
## V_z            0.47041    0.20691   2.273 0.022999 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Log-likelihood: -384.1
## BIC = 3664

```

```

cbind(True=c(beta=beta, alpha=alpha),
      Est=coef(mod))

```

##		True	Est
##	beta1	-0.5	-0.50813933
##	beta2	-0.5	-0.57343540
##	alpha1	0.0	0.02131236
##	alpha2	0.5	0.47040670